



WEBSITE BUILD BRIEF - VOLUME 2

Ranking, Content & Automation

How bold-clicks.com wins Austin, then Texas, then the surrounding states - and keeps winning without manual upkeep.

Prepared for: Bold Clicks (Stoneridge family)

Intended reader: Claude Code (build agent) and the site owner

Companion to: Bold Clicks Website Master Build Brief (Volume 1)

Date: September 18, 2026

WHAT THIS VOLUME ADDS

Geographic ranking strategy | Full keyword and content map | 12-month editorial calendar | Self-running monitoring and alerting | Scheduled automation jobs | Technical stack, design tokens, page copy | Lead routing, analytics and CRM spec | CI/QA gates | Launch plan | Claude Code operating instructions

00 / READ THIS FIRST

How to use Volume 2

This document extends the Master Website Build Brief. It does not replace it.

Volume 1 defined **what** Bold Clicks is, the brand system, the launch sitemap, the homepage blueprint, the service-page template, the technical SEO baseline, and the launch acceptance checklist. Volume 2 answers the questions Volume 1 left open: **how the site actually ranks** in Austin, across Texas, and in the surrounding states; **what content gets produced, when, and by whom**; and **how the whole system runs itself** so the owner never has to wonder whether tracking, indexing, performance, or forms are working.

PRECEDENCE RULES

1. If Volume 1 and Volume 2 conflict on a brand, logo, or 'do not fabricate' rule, Volume 1 wins.
2. If they conflict on SEO strategy, automation, content cadence, or implementation detail, Volume 2 wins (it is the more specific document).
3. Anything marked [OWNER] in either volume is a required input from the owner. Claude Code must create a clearly labeled placeholder in the content model and list it in the final hand-off report - never invent it, never ship it as visible copy.
4. Automation may draft, test, measure, and propose. Automation never publishes public content or changes tracking without an approval step.

The goal, stated plainly

Bold Clicks intends to be the number-one paid-media agency in the country. Nobody ranks nationally on day one, and a site that tries to rank everywhere at once usually ranks nowhere. So the ranking strategy is sequenced: **own Austin** first (where the operators, the proof, and the business address are), **then Texas** (Dallas-Fort Worth, Houston, San Antonio), **then the surrounding states** (Oklahoma, Louisiana, Arkansas, New Mexico), and only then compete on national head terms. Each tier unlocks when the previous tier produces measurable proof - rankings, leads, and case studies - that can be reused in the next tier.

Tier	Geography	Primary lever	Unlock condition
Tier 1	Austin metro (Austin, Round Rock, Cedar Park, Georgetown, Pflugerville, Leander, Lakeway, San Marcos, Kyle, Buda)	Google Business Profile + one strong Austin page + local citations + reviews + Austin-specific insights	Site live, GBP verified, tracking verified
Tier 2	Texas metros (Dallas-Fort Worth, Houston, San Antonio) and statewide 'Texas' modifiers	Texas hub page + industry pages with Texas proof + case studies with Texas clients + state-level content	At least 3 verified case studies, top-10 Austin positions on core terms for 60+ days
Tier 3	Oklahoma, Louisiana, Arkansas, New Mexico	Regional case studies, industry playbooks, benchmark reports, PR/links from regional trade groups	Tier 2 sustained, 6+ case studies, referring domains growing month over month
Tier 4	National head terms ('PPC agency', 'Google Ads agency')	Original data studies, tools, named experts, authority links, brand search volume	Tier 3 sustained and brand search demand visible in Search Console

Document map

Section	What it covers	Who acts on it
18	Geographic ranking strategy: Austin, Texas, surrounding states	Claude Code (pages), Owner (GBP, citations, reviews)
19	Automated competitive intelligence protocol	Claude Code (scripts), runs unattended
20	Full keyword and intent map by page cluster	Claude Code (page targeting), automation (rank tracking)
21	12-month editorial calendar	Automation drafts, Owner approves
22	Content refresh and freshness system	Automation detects, Owner approves
23	Automation architecture: the site runs itself	Claude Code builds it
24	Monitoring and alerting specification	Claude Code builds; alerts go to Owner
25	Scheduled jobs (cron) specification	Claude Code builds; GitHub Actions runs
26	AI content guardrails	Everyone
27	Owner reporting: weekly digest, monthly report, dashboard	Automation produces
28	Technical stack, repository layout, content schemas	Claude Code
29	Design system tokens and component specifications	Claude Code
30	Page copy drafts (homepage, free audit, Google Ads, Austin, About)	Claude Code implements; Owner edits
31	Lead capture, routing, and CRM specification	Claude Code + Owner (GoHighLevel)
32	Analytics, tag manager, and attribution implementation	Claude Code + Owner (account IDs)
33	CI/QA pipeline and pre-launch crawl	Claude Code
34	Launch plan and post-launch 30/60/90	Claude Code + Owner
35	Owner-supplied inputs checklist	Owner
36	Claude Code operating instructions (CLAUDE.md, commands, definition of done)	Claude Code
37	Risk register and anti-patterns	Everyone
A-C	Appendices: geography reference, event/schema reference, KPI glossary	Reference

18 / GEOGRAPHIC RANKING STRATEGY

Own Austin. Then own Texas.

Local and regional search for agency services is won with a small number of strong pages plus off-site signals - not with city-page sprawl.

Why geography is the first fight

When someone in Texas searches 'PPC agency' or 'Google Ads management,' Google usually localizes the results: a map pack of nearby agencies, then organic results that lean toward businesses with local relevance. National agencies with huge authority still show up, but a local specialist with a verified Google Business Profile, consistent citations, real reviews, and a page that speaks to the local market can outrank them for the localized version of the query. That is the opening. Bold Clicks is in Austin, which is also one of the most competitive local markets in the country for home services, med spas, dental, and legal - exactly the verticals Bold Clicks serves.

The three local ranking systems to win

System	What Google looks at	What Bold Clicks does
Map pack (Google Business Profile)	Relevance of categories/services, distance from searcher, prominence (reviews, citations, links, brand mentions)	Verified GBP with correct primary category, complete services, weekly posts, review flow from real clients, Q&A seeded with real buyer questions
Localized organic results	Location relevance on the page, entity consistency (name, address, phone), local links and mentions, topical authority	One authoritative Austin page, Austin-specific insights content, Organization schema with real address (if public), consistent NAP across the web
Regional / statewide organic	Topical authority, case studies and proof tied to the region, links from regional sources	Texas hub page, industry pages that reference Texas market conditions with real data, case studies naming Texas metros (with permission)

Location page rules (the anti-doorway standard)

Volume 1 already warned against mass city pages. This is the positive version of that rule: build very few location pages and make each one the best page on the internet for that market.

Route	Build when	Must contain (all required)
/austin-ppc-agency/	At launch	Austin-specific market commentary (CPC ranges by vertical from Bold Clicks' own accounts once available, seasonality, competitive density); map + service radius; Austin case studies or, pre-proof, Austin account teardown lessons; local FAQs (minimum spend in Austin, which platforms work for Austin home services, etc.); embedded GBP link; real address if public; named operator; internal links to every service page
/texas-ppc-agency/	After Tier 1 unlock (see section 00)	Statewide framing: DFW vs Houston vs San Antonio vs Austin differences in CPC, competition, and lead value; Texas case studies; Texas-specific compliance notes for legal and healthcare advertising; TDLR/state-licensing language for home services where relevant; links to industry pages
/dallas-ppc-agency/, /houston-ppc-agency/, /san-antonio-ppc-agency/	Only after a verified client in that metro exists	Same standard as Austin. If the page cannot cite a real client, real spend range, or real metro-specific data, do not publish it. A page that only swaps the city name is a doorway page and must be rejected in code review.

Route	Build when	Must contain (all required)
Surrounding-state pages	Never at launch. Revisit at Tier 3	Regional demand is served through industry pages, case studies, and benchmark content that mention the state naturally. A state page exists only if the team can say something genuinely useful about advertising in that state.

DOORWAY-PAGE REJECTION TEST (run in code review and in the monthly SEO audit)

A location page fails if any of these is true: (a) more than 60% of its body text is shared with another location page; (b) it contains no figure, screenshot, client reference, or observation that could only be true for that market; (c) it exists only to funnel to the same form with no unique buyer questions answered. Failing pages get noindex and a rewrite ticket - they do not get published.

Google Business Profile playbook [OWNER sets up; automation maintains]

- **Primary category:** Advertising agency. Secondary: Marketing agency, Internet marketing service. Do not add unrelated categories to chase volume.
- **Name:** Bold Clicks - exactly the legal/brand name, no keywords stuffed into the name (this is a suspension risk).
- **Address vs service area:** If Bold Clicks has a real, staffed Austin address, list it. If not, set a service-area business covering the Austin metro and hide the address. Never use a virtual office or co-working mailbox as a listed address.
- **Services list:** Google Ads management, Meta Ads management, Microsoft Ads management, YouTube Ads management, Conversion tracking and attribution, Free ad-account audit. Each with a short real description.
- **Description:** Reuse the positioning line from Volume 1. Mention Austin, Texas, and the platforms managed. No superlatives that cannot be proven.
- **Reviews:** Every closed client enters a GoHighLevel review-request workflow at day 45 and day 120. Respond to every review within 48 hours (automation drafts, owner sends).
- **Posts:** One post per week, drafted automatically from newly published Insights, case studies, or platform-update explainers. Owner approves in the weekly digest.
- **Q&A:** Seed 8-10 real buyer questions (minimum spend, contract length, account ownership, reporting cadence, which industries) with honest answers.
- **Photos:** Real team, real office/workspace, real screenshots (redacted). Refresh quarterly.
- **UTM the website link:** ?utm_source=google&utm_medium=organic&utm_campaign=gbp so GBP traffic is attributable.

Citation and profile consistency

Consistency of name, address, and phone (NAP) across the web is a prominence signal and a trust signal. Build these once, record them in the repo (content/citations.json), and let automation check them quarterly for drift.

Priority	Where	Notes
Core	Google Business Profile, Bing Places, Apple Business Connect, Yelp, BBB, LinkedIn Company Page, Facebook Page, Instagram	Identical NAP, identical logo asset, identical description, tracked website link
Agency directories	Clutch, UpCity, DesignRush, GoodFirms, Sortlist, The Manifest, Agency Spotter	Complete profiles; collect verified client reviews on Clutch and UpCity specifically - they rank for 'best PPC agency austin' queries themselves

Priority	Where	Notes
Platform directories	Google Partners directory (only once Partner status is actually earned), Meta Business Partners (only if earned), Microsoft Advertising Partners (only if earned)	Never claim status that is not current. Automation checks the badge status quarterly and flags expiry.
Local	Austin Chamber of Commerce, Greater Austin Black Chamber / Hispanic Chamber (if applicable), Austin Business Journal Book of Lists submissions, Texas Association of Business, local networking groups	Chamber memberships typically include a directory link - these are legitimate local links
Vertical	Dental/ortho practice-management vendor partner pages, med-spa software partner directories, legal marketing directories, home-services franchise vendor lists	Only where Bold Clicks has a genuine relationship or client; these double as industry-relevant links

Austin-specific content angles (feed these into the calendar in section 21)

- What Google Ads actually costs in Austin by vertical - built from Bold Clicks account data with methodology disclosed; until then, framed as 'what we see' with ranges, never fake precision.
- Austin home-services seasonality: HVAC and roofing demand curves, storm-season spikes, and how bids should move with them.
- Med spa advertising in Austin: creative fatigue in a saturated market and what actually differentiates offers.
- Austin dental and orthodontic paid search: emergency vs elective intent, the Round Rock / Cedar Park suburban competition pattern.
- Legal PPC in Travis County: case-type CPC reality, intake-quality problems, and call-tracking lessons.
- Landing-page teardowns of anonymized Austin businesses (with permission or fully anonymized).
- Austin event-driven demand (SXSW, ACL, F1 weekend) and whether it matters for local service advertisers - an honest analysis, not hype.

Texas and surrounding-state expansion signals

Rather than building state pages, the site accumulates regional relevance through content and proof. Automation tracks Search Console queries containing each metro or state name and reports them monthly so the owner can see which regions are already generating demand. When a region shows organic impressions plus at least one signed client, that region is a candidate for its own page (section 18, location page rules).

Region	Demand signal to watch	Content lever before a page exists
Dallas-Fort Worth	GSC queries containing dallas, fort worth, dfw, plano, frisco, arlington	Industry pages reference DFW competition; case studies with DFW clients
Houston	houston, the woodlands, katy, sugar land, pearland	Home-services and legal content referencing Houston CPC and storm seasonality
San Antonio	san antonio, new braunfels, boerne	Healthcare and dental content; proximity to Austin makes shared case studies natural
Oklahoma	oklahoma city, tulsa, okc	Benchmark reports segmented by state; regional trade-group contributions
Louisiana	new orleans, baton rouge, lafayette, shreveport	Home-services (roofing, storm restoration) playbooks
Arkansas	little rock, fayetteville, bentonville	Med-spa and dental playbooks; Northwest Arkansas growth-market commentary

Region	Demand signal to watch	Content lever before a page exists
New Mexico	albuquerque, santa fe, las cruces	Healthcare paid-media content; bilingual campaign lessons where real

19 / COMPETITIVE INTELLIGENCE PROTOCOL

Never guess who you are competing with

The competitive set changes as the site moves through tiers. Automation keeps the list current; humans decide what to do about it.

Volume 1 named four national agencies as pattern references. Those are not the competitors Bold Clicks meets on an Austin SERP. The Austin and Texas competitive set will include local PPC shops, full-service agencies with strong local SEO, and directory sites (Clutch, UpCity, Yelp) that rank for 'best PPC agency austin' themselves. This section defines a script, not a static list, so the competitive picture is always current and never based on stale assumptions.

Monthly SERP scan (automated)

1. For every query in the tracked set (section 20, priority P1 and P2, with and without geographic modifiers), fetch the top 20 organic results and the map-pack entries using a SERP data provider [OWNER chooses: DataForSEO, SerpApi, or Semrush API - the repo must support any of them behind one adapter].
2. Store results in data/serp/YYYY-MM/ as JSON, committed to the repo so history is diffable.
3. Classify each result domain: local agency, national agency, directory/aggregator, platform doc (Google), publisher, other.
4. Compute share of voice per query group and per geography tier.
5. For any domain appearing in the top 10 for three or more P1 queries, run a page-level profile: title, H1, word count, presence of case studies, offer type, form fields, schema types, estimated Core Web Vitals via the PageSpeed Insights API.
6. Produce docs/competitive/YYYY-MM.md: what changed since last month, new entrants, who lost positions, and which competitor pages Bold Clicks' equivalents are weaker than (and on which dimension).
7. Open a GitHub issue labeled 'competitive-gap' for each concrete, actionable gap. Include the evidence.

What to capture per competitor page

Dimension	Why it matters	Bold Clicks response rule
Offer	Free audit vs consultation vs pricing page - the offer sets conversion expectations for the SERP	Match or beat clarity; never match with a fake offer
Proof density	Number of case studies with numbers, named clients, review count and rating	Gap = case study production priority (section 21)
Topical depth	Do they have platform pages, industry pages, calculators, benchmark content?	Prioritize the missing asset type that section 20 already planned
Local signals	Address, service-area language, GBP presence and review count	Feed into GBP review-request cadence
Page experience	PSI mobile score, LCP/INP/CLS field data	Bold Clicks must stay in the 'good' band; a competitor with faster pages is a red flag
Links	Referring domains count and notable sources (via the chosen API)	Informs link-worthy asset priorities (section 13A in Volume 1)

RULE

Competitive intelligence informs priorities. It never licenses copying design, copy, or structure. Any output from this protocol that reads like 'replicate competitor X' is wrong; the right output is 'competitor X proves buyers want Y; build Bold Clicks' original version of Y.'

20 / KEYWORD AND INTENT MAP

One page, one primary intent, unlimited supporting depth

Priorities: P1 = launch, P2 = first 90 days, P3 = months 4-12. Volumes are not listed because they must be pulled live at build time and stored in data/keywords.json.

KEYWORD DATA RULE

Claude Code must pull current volume, CPC, and difficulty from the chosen keyword API at build time and store them in data/keywords.json with a pulled_at date. This brief supplies the map (which terms, which page, which intent, which priority). It deliberately does not supply numbers that would be stale within weeks.

Cluster 1: Core agency terms

Primary term	Supporting terms	Target page	Intent	Priorit y
ppc agency austin	austin ppc agency, ppc management austin tx, paid search agency austin	/austin-ppc-agency/	Local commercial	P1
google ads agency austin	austin google ads management, google adwords agency austin	/austin-ppc-agency/ (H2 section) + /google-ads-management/	Local commercial	P1
paid media agency	paid advertising agency, paid media management, performance marketing agency	/ (homepage)	Commercial, broad	P1
ppc management services	ppc management company, pay per click management	/services/	Commercial	P1
ppc agency texas	texas ppc agency, google ads agency texas, paid search texas	/texas-ppc-agency/	Regional commercial	P2
ppc agency dallas / houston / san antonio	google ads management dallas, etc.	Metro pages (only when unlocked)	Local commercial	P3
best ppc agency	top ppc agencies, best google ads agencies	Case studies hub + directory profiles (Clutch/UpCity)	Commercial investigation	P3

Cluster 2: Platform terms

Primary term	Supporting terms	Target page	Intent	Priority
google ads management	google ads management services, google ads management company, adwords management	/google-ads-management/	Commercial	P1
performance max management	pmax agency, performance max vs search	/google-ads-management/ (section) + Insight article	Commercial + informational	P2
meta ads management	facebook ads management, instagram ads agency, facebook ads agency, paid social agency	/meta-ads-management/	Commercial	P1
microsoft ads management	bing ads management, bing ads agency, microsoft advertising agency	/microsoft-ads-management/	Commercial	P2
youtube ads management	youtube advertising agency, youtube ads agency, video ads management	/youtube-ads-management/	Commercial	P2
google local services ads management	lsa management, local services ads agency	Home-services industry page section; standalone page only if demand is proven	Commercial	P3

Cluster 3: Measurement terms

Primary term	Supporting terms	Target page	Intent	Priority
ppc conversion tracking	google ads conversion tracking setup, conversion tracking agency	/conversion-tracking/	Commercial + informational	P1
offline conversion tracking google ads	enhanced conversions for leads, google ads crm integration, data manager google ads	/conversion-tracking/ + Insight guide	Informational (high intent)	P2
call tracking for google ads	ppc call tracking, callrail google ads	Insight guide -> links to /conversion-tracking/	Informational	P2
how to track qualified leads	lead quality tracking, qualified lead vs lead	Insight guide	Informational	P2
gohighlevel google ads integration	ghl google ads conversion, gohighlevel attribution	Insight guide (Bold Clicks has first-hand experience here)	Informational	P3

Cluster 4: Industry terms

Primary term	Supporting terms	Target page	Intent	Priority
dental ppc agency	google ads for dentists, dental google ads management, dental ppc	/industries/dental/	Commercial	P2
orthodontist ppc	orthodontic marketing google ads, invisalign google ads	/industries/orthodontics/	Commercial	P3
podiatry marketing google ads	podiatrist ppc, foot and ankle google ads	/industries/podiatry/	Commercial	P3
ent marketing ppc	sinus google ads, ent google ads	/industries/ent-sinus/	Commercial	P3
primary care google ads	medical practice ppc, doctor ppc agency	/industries/primary-care/	Commercial	P3
med spa ppc agency	med spa google ads, medical spa facebook ads, aesthetics ppc	/industries/med-spa/	Commercial	P2
legal ppc agency	law firm google ads, attorney ppc, personal injury google ads	/industries/legal/	Commercial	P2
home services ppc	hvac google ads, roofing ppc, plumber google ads, roofing google ads agency	/industries/home-services/	Commercial	P1 (Stoneridge has roofing proof)

Cluster 5: Educational / buyer-question terms

Primary term	Format	Supports	Priority
how much does ppc management cost	Pricing-model explainer with honest ranges and what changes the price	/services/, /free-ad-audit/	P1
google ads budget for small business	Budget planning guide + break-even calculator	/google-ads-management/	P2
why are my google ads leads low quality	Diagnostic article with search-term, landing-page, and tracking causes	/conversion-tracking/	P1
performance max vs search campaigns	Decision framework with account examples	/google-ads-management/	P2
what should a ppc agency report	Reporting standard article (Bold Clicks' own template shown)	/about/, /services/	P2
google ads account audit checklist	Audit framework (the same one the free audit uses)	/free-ad-audit/	P1
cost per lead benchmarks by industry	Benchmark report (only when sample size and methodology are real)	Industry pages	P3
facebook ads vs google ads for leads	Platform-fit article	/meta-ads-management/	P2

Cluster 6: Brand and navigational

'bold clicks', 'bold clicks agency', 'bold clicks austin', 'bold clicks reviews'. Homepage, About, and Case Studies own these. Automation tracks brand-query impressions monthly as the leading indicator of Tier 4 readiness. Register bold-clicks and boldclicks handles on every major platform at launch even if unused, to protect the brand SERP.

Keyword-to-page assignment rules

- Each commercial page has exactly one primary term in the title tag and H1 family. Supporting terms appear in H2s, body, FAQs, and internal anchors - naturally.
- Geographic modifiers live on location pages. Service pages mention Austin and Texas in copy and schema but do not chase 'austin' in their titles; that would cannibalize the Austin page.
- Educational articles never target a term already owned by a commercial page. If an article starts ranking for a commercial term, the automation flags it and the fix is an internal link plus a canonical decision - not a second commercial page.
- Rank tracking is by query group (cluster) and by page, not by a single vanity keyword.

21 / 12-MONTH EDITORIAL CALENDAR

Fewer pieces, each one earning its place

Cadence: 2 substantial Insights per month, 1 case study per month once clients exist, 1 platform-update explainer whenever Google or Meta ships a material change, 1 quarterly benchmark or data asset.

Month 1 begins at launch. Every item lists the target query it supports and the page it feeds internal links to. Automation drafts the first version from the outline plus any Bold Clicks account data the owner has approved for use; the owner or a named operator adds first-hand experience and approves. Nothing on this calendar publishes without that approval (section 26).

Mo	Piece	Target query	Feeds	Format
1	Google Ads Account Audit Checklist (the Bold Clicks framework)	google ads account audit checklist	/free-ad-audit/	Framework + downloadable
1	Why Your Google Ads Leads Are Low Quality (and the 7 places we look first)	why are my google ads leads low quality	/conversion-tracking/	Diagnostic
2	What PPC Management Costs in 2026: pricing models, what changes the price, what to avoid	how much does ppc management cost	/services/	Pricing explainer
2	Austin Home-Services Paid Search: seasonality, storm spikes, and bid timing	roofing google ads agency austin / hvac google ads austin	/industries/home-services /, /austin-ppc-agency/	Local playbook
3	Performance Max vs Search: a decision framework from real accounts	performance max vs search campaigns	/google-ads-management /	Decision guide
3	How We Track Qualified Leads (GoHighLevel + Google Ads offline conversions walkthrough)	how to track qualified leads / gohighlevel google ads integration	/conversion-tracking/	Technical guide
3	Case Study #1 [OWNER: first verified client]	brand + industry term	/case-studies/	Case study
4	Break-even CPL / CAC Calculator + methodology page	google ads budget for small business	/google-ads-management /	Interactive tool
4	Med Spa Advertising: creative fatigue and offer testing in saturated markets	med spa ppc agency / med spa facebook ads	/industries/med-spa/	Vertical playbook
4	Case Study #2	vertical term	/case-studies/	Case study
5	Facebook Ads vs Google Ads for Lead Gen: when each wins	facebook ads vs google ads for leads	/meta-ads-management/	Platform-fit guide
5	Legal PPC: case-type economics, intake quality, and call tracking	legal ppc agency / law firm google ads	/industries/legal/	Vertical playbook
5	Case Study #3	vertical term	/case-studies/	Case study
6	What a PPC Agency Should Report Every Month (our template, explained)	what should a ppc agency report	/about/, /services/	Standards piece
6	Dental and Orthodontic Paid Search: emergency vs elective intent	dental ppc agency / google ads for dentists	/industries/dental/, /industries/orthodontics/	Vertical playbook
6	Quarterly data asset: Austin CPC ranges by vertical (only if sample size supports it)	google ads cost austin	/austin-ppc-agency/	Benchmark

Mo	Piece	Target query	Feeds	Format
7	Search-Term Waste Checklist: 25 negative-keyword patterns we add on day one	google ads negative keywords list	/google-ads-management /	Checklist
7	Microsoft Ads for Local Service Businesses: when it is worth it	microsoft ads management / bing ads agency	/microsoft-ads-managem ent/	Platform guide
7	Case Study #4	vertical term	/case-studies/	Case study
8	Landing-Page Teardown: three anonymized Texas service businesses	ppc landing page best practices	/services/	Teardown
8	Call Tracking for Google Ads: setup, pitfalls, and what counts as a real lead	call tracking for google ads	/conversion-tracking/	Technical guide
9	YouTube Ads for Lead Gen: audience, creative, and where it sits in the funnel	youtube ads management	/youtube-ads-manageme nt/	Platform guide
9	Texas Paid Media: DFW vs Houston vs San Antonio vs Austin - what actually differs	ppc agency texas	/texas-ppc-agency/ (if unlocked)	Regional analysis
9	Case Study #5	vertical term	/case-studies/	Case study
10	Google Ads for Podiatry / ENT / Primary Care: service-line economics	podiatry marketing google ads / ent marketing ppc	Healthcare industry pages	Vertical playbook
10	Quarterly data asset: cost-per-lead ranges by vertical (Texas)	cost per lead benchmarks by industry	Industry pages	Benchmark
11	Local Services Ads vs Search for Home Services	google local services ads management	/industries/home-services /	Platform guide
11	Enhanced Conversions and Data Manager: a practical setup guide	enhanced conversions for leads	/conversion-tracking/	Technical guide
11	Case Study #6	vertical term	/case-studies/	Case study
12	Year-One Review: what we changed across accounts and why (original data)	brand	/about/	Authority piece
12	Annual benchmark report (if data supports) + PR distribution	ppc benchmarks texas	Industry pages, /insights/	Link asset

Standing series (event-driven, not calendar-driven)

- **Platform update explainers.** Automation watches the Google Ads, Google Search Central, Meta for Business, and Microsoft Advertising blogs (RSS). When a post matches material-change keywords (bidding, attribution, conversion, Performance Max, privacy, policy, consent), it drafts an explainer with a mandatory 'what we are changing in client accounts' section that the owner must fill. Published only after approval.
- **Case studies.** Triggered by the CRM: when a client passes 90 days and the owner marks results as verified in GoHighLevel, automation drafts the case study from the template in Volume 1 section 06 and opens a review PR.
- **Austin market notes.** Short monthly note on what changed in Austin CPCs and lead volume in Bold Clicks accounts. Only published if the owner confirms the numbers.

22 / CONTENT REFRESH AND FRESHNESS SYSTEM

Content decays. The system notices before the owner does.

Every indexable page carries a review cadence and a set of decay triggers. Automation opens a refresh task when either fires.

Page type	Scheduled review	Decay triggers (any one opens a task)	What a refresh includes
Homepage	Quarterly	Clicks down 20%+ over 8 weeks vs prior 8; CTR down 25%+; new case study published; positioning line changed	Proof strip update, case-study cards, FAQ additions from real sales calls, CTA copy test
Platform service pages	Quarterly	Position for primary term drops 3+ places for 4 weeks; platform ships a material feature change; competitor page outranks on the same term	Deliverables list, process section, new FAQs, platform-specific screenshots, internal links to new insights
Industry pages	Semi-annually	New vertical case study; benchmark asset updated; impressions rising with no click growth	Add the case study, update KPI ranges, add objections heard on sales calls
Location pages	Quarterly	GBP insights shift; local competitor enters top 3; new local client	Market commentary, local proof, local FAQ
Insights articles	Annually, or on trigger	Clicks down 30%+ over 12 weeks; the referenced platform feature changes; broken external link; article age > 18 months	Update facts, add new first-hand data, fix links, update 'last reviewed' date only if material change was made
Case studies	On trigger	Client relationship ends; results materially change; permission revoked	Update, anonymize, or unpublish (with 410 or 301 to the hub) - never leave stale claims live
Calculators / tools	Quarterly	Input assumptions (industry defaults) older than 6 months; error rate in logs	Recalibrate defaults with methodology note
Legal pages	Annually and on tracking change	Any new tag, cookie, or vendor added	Privacy policy vendor list and consent copy

Decay detection (automated, weekly)

1. Pull Search Console performance by page and by query group for the trailing 16 weeks (Search Console API).
2. Compare the most recent 8 weeks with the prior 8 weeks for clicks, impressions, CTR, and average position.
3. Apply the trigger thresholds above per page type. Ignore pages with fewer than 50 impressions in the window (too noisy).
4. For triggered pages, gather context: top losing queries, competitor pages now above (from the section 19 SERP data), and any recent commits touching the page.
5. Open a GitHub issue labeled 'content-refresh' with the evidence and a proposed refresh scope. If the page is an Insight article, also draft the refreshed version in a branch and attach a preview link.
6. Summarize all triggered pages in the weekly digest (section 27).

FRESHNESS HONESTY RULE

Do not bump 'updated' dates, sitemap lastmod, or Article dateModified unless the content materially changed. Fake freshness is detectable and erodes trust with both Google and buyers. The CI pipeline blocks a lastmod change on any file whose content diff is under 5% unless the commit is tagged 'material-update'.

Pruning rules

- A page with zero clicks and under 100 impressions for 6 months is a candidate for consolidation (301 to the strongest related page) or improvement. Automation lists candidates quarterly; the owner decides.
- Never delete a page that has external links pointing to it without a 301 to the most relevant replacement.
- Thin drafts that were never approved are deleted from the branch after 90 days so the backlog stays honest.

23 / AUTOMATION ARCHITECTURE

The site runs itself. The owner approves, not operates.

Design principle: automate observation, drafting, testing, and reporting completely; keep publishing and tracking changes behind a one-click human approval.

Operating principles

1. **One inbox, one digest.** The owner receives one weekly digest email/SMS (via GoHighLevel) that lists everything needing a decision, plus immediate alerts only for outages, broken forms, or tracking failures. Nothing else pings the owner.
2. **Every job is idempotent and logged.** A job can run twice without doing damage, and every run writes a JSON log to `data/runs/` committed to the repo, so 'is it working?' is answerable from Git history.
3. **Automation proposes through pull requests.** Content drafts, refresh rewrites, meta-description fixes, internal-link suggestions, and schema corrections arrive as PRs with a preview deployment. Approving the PR is the publish action.
4. **Guardrails are code, not policy.** CI blocks fabricated numbers (see section 26 linter), placeholder text, lastmod abuse, missing metadata, and accessibility regressions. If it is important, it is a failing check, not a reminder.
5. **Fail loudly, degrade gracefully.** If a data provider is down, the job reports 'skipped: provider unavailable' rather than silently producing an empty report that looks like good news.
6. **No silent dependencies on one person.** All credentials live in GitHub Actions secrets or the hosting provider's environment; the README documents how to rotate every one of them.

Layers

Layer	Responsibility	Implemented with
Hosting + deploy	Static/SSR build, preview deployments per PR, atomic production deploys, instant rollback	Vercel (recommended) or Cloudflare Pages; GitHub as source of truth
Scheduling	Runs every job on a cron; retries; secrets	GitHub Actions workflows (schedule: cron); optionally a lightweight queue on Cloudflare Workers if jobs exceed Actions limits
Data collection	Search Console, GA4, PageSpeed Insights / CrUX, uptime, SERP provider, RSS, GBP performance (via API where available)	Node/TypeScript scripts in /automation with a provider adapter pattern; results written to /data as JSON
Analysis	Decay detection, competitive diffs, anomaly detection, freshness checks, keyword cannibalization detection	TypeScript with simple statistical rules (no ML needed); Claude Code invoked via the Claude Code GitHub Action for narrative analysis and drafting
Generation	Content drafts, refresh rewrites, GBP posts, weekly digest narrative, meta fixes	Claude Code GitHub Action with the repo's CLAUDE.md, brand voice file, and 'no fabrication' linter
QA	Lint, typecheck, build, Lighthouse CI, axe, link check, schema validation, metadata uniqueness, no-fabrication lint	GitHub Actions on every PR; required checks before merge
Approval + publish	PR review; merge deploys production	GitHub PR with preview link in the digest
Reporting	Weekly digest, monthly report, always-on dashboard	GoHighLevel email/SMS for digest; Looker Studio dashboard on GA4 + GSC + a BigQuery or Google Sheet fed by /data exports

Data flow (read left to right)



What the owner actually touches

Moment	Owner action	Time
Weekly (Monday 7am CT digest)	Read digest; approve or reject listed PRs by clicking merge/close; answer any [OWNER] questions	10-15 minutes
Immediate alert (rare)	Read the alert; if it is a site outage, tracking failure, or broken form, the runbook link in the alert says exactly what to do	As needed
Monthly	Read the monthly report; confirm any numbers automation wants to publish (benchmarks, case-study figures)	20 minutes
Quarterly	Review pruning candidates, citation drift report, partner-badge status, and the competitive summary	30 minutes

DEFINITION OF 'AUTOMATED'

A task is automated only when it runs on a schedule without a human starting it, writes a log the owner can find, reports its own failure, and either takes a safe action or opens a clearly labeled PR/issue. A script that exists but must be remembered is not automation.

24 / MONITORING AND ALERTING SPECIFICATION

'Is it working?' must always have an answer

Each monitor lists the tool, cadence, threshold, automatic action, and where the alert goes. Severity: S1 = immediate SMS + email; S2 = same-day email; S3 = weekly digest only.

Monitor	Tool	Cadence	Threshold	Automatic action / alert	Sev
Uptime (homepage, /free-ad-audit/, /google-ads-management/)	External HTTP probe from 2 regions (Better Uptime / UptimeRobot / Checkly [OWNER picks])	Every 1 min	Non-200 or > 5s for 2 consecutive checks	Page owner via SMS; auto-create incident issue	S1
SSL certificate	Same probe	Daily	Expiry < 14 days or handshake failure	Issue + email	S1 if failed, S2 if expiring
Form synthetic test	Playwright job submits the audit form with a flagged test email (test+synthetic@) and verifies (a) 200, (b) thank-you state, (c) generate_lead event fired in GA4 DebugView/measurement protocol, (d) contact appears in GoHighLevel via API with UTMs intact	Every 6 hours	Any step fails	SMS + issue with the failing step	S1
Phone link and booking link	Playwright checks href, tel: format, and that the Calendly/GHL calendar loads	Daily	Broken or not loading	Issue + email	S2
Core Web Vitals (field)	CrUX API for origin and top 10 URLs; PSI API lab run on mobile	Daily (CrUX), weekly (PSI)	Any metric leaves 'good' at p75 for 7 days; or PSI mobile performance < 85	Issue with the failing URL and metric; digest	S2
Indexing and coverage	Search Console URL Inspection API on all canonical URLs; sitemap status	Weekly	Any canonical URL not indexed for 14 days post-publish; any 'noindex detected' on a page that should be indexed; sitemap errors	Issue; digest	S2
Accidental noindex / robots block	Crawl of production comparing robots meta and X-Robots-Tag against the canonical route list	Daily	Mismatch	SMS (this is a revenue-killing failure)	S1
Rankings by query group	Search Console average position and clicks per cluster; SERP provider for map pack	Weekly	Cluster position worsens 3+ places for 2 weeks; map-pack drop	Digest with diagnosis	S3
Broken links (internal + external)	Crawler script on the sitemap	Weekly	Any 4xx/5xx internal; > 3 external	Auto-PR fixing internal; issue for external	S3

Monitor	Tool	Cadence	Threshold	Automatic action / alert	Sev
Redirect health	Test the redirect map file against production	Weekly	Chains > 1 hop, loops, or 302 where 301 expected	Auto-PR	S3
404 spike	Server logs / Vercel analytics	Daily	> 50 unique 404 paths/day or a 3x increase	Issue with top paths	S2
Tracking integrity	GTM container export diff + GA4 realtime check that key events (cta_click, form_start, generate_lead, phone_click) fired in the last 24h	Daily	Container changed without a matching PR; or zero events for 24h on a page with traffic	SMS for zero-events; issue for container drift	S1 / S2
Consent behavior	Playwright runs with consent denied and asserts no ad/analytics cookies set before consent	Weekly	Any violation	Issue	S2
Schema validity	Validate all JSON-LD against schema.org types and Google's structured-data rules (via the Rich Results Test API where available or a validator library)	On every PR + weekly	Any error	Block PR / issue	S2
Metadata uniqueness	Script over the built HTML: unique title, description, H1, canonical present, OG tags present	On every PR	Duplicates or missing	Block PR	S3
Citation drift	Fetch each core citation and compare NAP to content/citations.json	Quarterly	Mismatch	Issue with correction instructions	S3
Partner badge status	Owner-confirmed field in content/proof.json with an expiry date	Monthly	Expiry within 30 days	Digest	S3
Dependency security	GitHub Dependabot + npm audit	Daily	High/critical vulnerability	Auto-PR; issue	S2
Backups	Repo is the backup; additionally export GHL contacts and GA4 config monthly to a private bucket	Monthly	Export failed	Issue	S2
Domain + DNS	WHOIS expiry check; DNS record diff vs infra/dns.json	Weekly	Expiry < 60 days; record drift	SMS if < 30 days; issue otherwise	S1 / S2
GBP anomalies	GBP Performance API (calls, direction requests, website clicks) if available; otherwise monthly manual export	Weekly	Drop of 40%+ week over week; profile suspension	Digest; SMS on suspension	S3 / S1

Alert routing

- S1 alerts go through GoHighLevel to the owner's SMS and email simultaneously, and open a GitHub issue labeled 'incident' with a runbook link.

- S2 alerts open an issue and send one email the same day. Duplicate alerts for the same condition are suppressed for 24 hours.
- S3 items appear only in the weekly digest.
- Every alert includes: what broke, when it was first detected, what the automation already tried, and the exact next step for a human.

Runbooks (write these into /docs/runbooks/ at build time)

Runbook	Covers
site-down.md	Check hosting status page, roll back last deploy from the hosting dashboard, verify DNS, escalate to developer
form-broken.md	Check GHL webhook status, check hosting function logs, replay the synthetic test, temporary fallback to mailto capture
tracking-zero-events.md	Confirm GTM container published version, confirm consent script not blocking, check GA4 property/stream IDs, check CSP headers
noindex-detected.md	Identify the source (env var, header, meta), fix, redeploy, request re-indexing via Search Console
cwv-regression.md	Compare last 3 deploys, check new third-party scripts, check image sizes, run PSI with the failing URL
gbp-suspended.md	Common causes, reinstatement form steps, what evidence to gather

25 / SCHEDULED JOBS SPECIFICATION

Every job, its schedule, its output, and its approval gate

All times are America/Chicago. Implement each as a GitHub Actions workflow in `.github/workflows/` with a matching script in `/automation/`.

Job	Schedule	What it does	Output	Gate
nightly-collect	Daily 02:00	Pull GSC (pages, queries, coverage), GA4 (sessions, events, conversions by source), CrUX/PSI, uptime summary, 404s. Write to <code>/data</code> .	Commit to data branch; open issue on provider failure	None (data only)
daily-integrity	Daily 06:00	Noindex/robots check, tracking-zero-events check, SSL, DNS drift, dependency audit.	S1/S2 alerts as defined in section 24	None (alerts only)
form-synthetic	Every 6 hours	Playwright end-to-end submission through to GoHighLevel and GA4 event verification.	S1 alert on failure	None
weekly-seo-analysis	Monday 04:00	Decay detection (section 22), cannibalization detection, rank-cluster movement, broken links, redirect health, coverage.	Issues labeled content-refresh / seo-fix; auto-PRs for internal link and redirect fixes	PRs require merge
weekly-digest	Monday 07:00	Compile everything from the week into a plain-English digest: wins, problems, PRs awaiting approval, [OWNER] questions, numbers vs last week.	Sent via GoHighLevel email + short SMS summary; archived in <code>/docs/digests/</code>	None (read only)
content-draft	1st and 15th, 05:00	Take the next calendar item (section 21), assemble the outline, pull any approved account data from <code>/data/approved-facts.json</code> , draft via Claude Code Action, run the no-fabrication linter.	PR with preview deployment and a checklist of [OWNER] inputs still needed	PR merge = publish
platform-update-watch	Every 4 hours	Poll RSS feeds for Google Ads, Search Central, Meta for Business, Microsoft Advertising. Match against material-change keywords.	Draft explainer PR with the mandatory 'what we are changing' section left as [OWNER] placeholder	PR merge
case-study-trigger	Daily 05:30	Query GoHighLevel for clients tagged 'results-verified' without a matching case study.	Draft case-study PR from template	PR merge + client permission field set
gbp-post-draft	Thursday 05:00	Draft next week's GBP post from the most recent published Insight or case study.	Included in weekly digest for one-click approval; posted via API if available, else queued for owner	Approval
monthly-competitive	1st, 03:00	Section 19 SERP scan and competitor page profiling.	<code>docs/competitive/YYYY-M</code> M.md + competitive-gap issues	None
monthly-report	1st, 07:00	Full monthly report (section 27) with leads, booked calls, source attribution, rankings, CWV, content shipped, and next month's calendar.	PDF + Looker Studio refresh; sent via GHL	None
monthly-keyword-refresh	1st, 03:30	Re-pull volume/CPC/difficulty for the keyword map; flag new query opportunities from GSC (queries with impressions but no target page).	Updated <code>data/keywords.json</code> PR; opportunity issues	PR merge

Job	Schedule	What it does	Output	Gate
quarterly-citations	Jan/Apr/Jul/Oct 1st	Citation drift check; partner-badge expiry; pruning candidates; calculator default recalibration.	Issues + digest section	Owner decisions
quarterly-benchmark	Jan/Apr/Jul/Oct 5th	If approved-facts contains enough anonymized account data (minimum sample rule in section 26), rebuild benchmark tables and draft the update.	PR	PR merge + owner confirms numbers
backup-export	1st, 01:00	Export GHL contacts/pipelines, GA4 admin config, GTM container, DNS zone to private storage.	Issue on failure	None

Implementation notes for Claude Code

- Put all provider calls behind /automation/providers/* with a common interface (search console, analytics, vitals, serp, crm, notify). The owner may swap SERP or uptime vendors without touching job logic.
- Use a single /automation/config.ts for thresholds so the owner (or a future developer) can tune sensitivity in one place. Ship the thresholds from section 24 as defaults.
- Every job writes /data/runs//.json with status, duration, counts, and errors. The weekly digest includes a one-line 'all jobs ran' or lists the ones that did not.
- Prefer the Claude Code GitHub Action for drafting and narrative tasks; pass it CLAUDE.md, docs/voice.md, data/approved-facts.json, and the target outline. Never pass it credentials.
- GitHub Actions scheduled workflows can be delayed under load; do not depend on exact minute timing. Jobs must tolerate running late or twice.
- Keep third-party API usage within free/low tiers where possible and log cost estimates in the run log; the owner should never get a surprise bill.

26 / AI CONTENT GUARDRAILS

Automation writes drafts. Humans own claims.

The site's authority depends on original experience. These rules keep AI-assisted production from turning into a content farm.

Hard rules

1. No AI-drafted content publishes without a named human reviewer recorded in the content's frontmatter (reviewedBy, reviewedAt).
2. Every published Insight must contain at least one 'first-hand block': an observation, screenshot, number, or lesson from an actual Bold Clicks or Stoneridge account (anonymized as needed). The linter fails a draft that has none.
3. Numbers appear only if they exist in data/approved-facts.json (owner-confirmed) or are cited to a public source with a URL. Any other number is a linter failure.
4. Benchmarks require a disclosed methodology, minimum sample size (at least 10 accounts or 90 days of data per segment, whichever the owner sets), and a date range.
5. No draft may claim certifications, partner status, client names, awards, or results not present in content/proof.json.
6. Drafts follow docs/voice.md (the Volume 1 voice rules plus a banned-phrase list). The linter flags banned phrases; reviewers reject 'could be any agency' copy.
7. Publishing volume is capped in config (default: 4 Insights per month). Exceeding the cap requires an owner override, which is logged.

The no-fabrication linter (build this)

Check	Rule	On failure
Unsourced numbers	Any digit sequence with %, \$, x, ROAS, CPL, CAC, leads, or spend context must trace to approved-facts.json (by fact id in frontmatter) or a citation URL	Block
Proof claims	Words like 'certified', 'partner', 'award', 'clients include', 'ranked #1' must match proof.json entries	Block
Placeholders	[OWNER], TODO, lorem, TBD, XX in any publishable file	Block
First-hand block	Insight articles must contain a component with non-empty content	Block
Banned phrases	docs/voice.md banned list (unlock your potential, cutting-edge, 360-degree, results-driven without proof, synergy, game-changing, etc.)	Warn (reviewer decides)
Duplication	Any 40-word window shared with another page on the site	Warn; block if > 30% of the page
Lastmod abuse	lastmod/dateModified changed with < 5% content diff and no material-update tag	Block
Metadata	Title 45-65 chars, description 120-160 chars, unique across the site, H1 present and single	Block

WHY THIS MATTERS FOR RANKING

Google's scaled-content and spam policies target unoriginal content produced primarily for rankings regardless of how it was made. The defense is not 'hide that AI helped'; it is publishing content that contains things only Bold Clicks could know. The first-hand block requirement is the single most important ranking control in this document.

27 / OWNER REPORTING

One weekly digest, one monthly report, one dashboard

The owner should be able to answer 'is the site working and is it growing?' in under a minute, any day.

Weekly digest (Monday 07:00 CT, email + SMS summary)

Block	Content
Status line	'All systems normal' or the list of anything that is not, in one sentence each
Leads	Form submissions, booked calls, phone clicks - this week vs last week vs 4-week average; by source (organic, paid, GBP, direct, referral)
Needs your decision	PRs awaiting approval with preview links and a one-line description each; [OWNER] questions from drafts; GBP post to approve
Rankings	Movement by cluster (core, platform, measurement, industry, local) with the top 3 gainers and losers; map-pack position for tracked Austin terms
Site health	CWV status, indexing coverage, any refresh tasks opened
Shipped this week	Merged PRs, published pages
Coming next	Next calendar item and its due date

Monthly report (1st of month, PDF in Bold Clicks branding)

- Business outcomes first: leads, qualified leads, booked calls, and closed clients attributed to the site (from GoHighLevel pipeline stages), with source and campaign attribution.
- Organic growth: clicks, impressions, CTR, average position by cluster and by geography tier; brand-query trend; pages entering top 10 / top 3.
- Local: GBP calls, direction requests, website clicks; review count and rating; map-pack positions.
- Technical health: CWV field data, coverage, errors resolved, uptime.
- Content: shipped, refreshed, pruned; calendar adherence; owner-input backlog.
- Authority: new referring domains, notable mentions, citation status.
- Competitive: two-paragraph summary from the monthly scan with the one gap that matters most.
- Next month: calendar, planned experiments, tier-unlock status against the criteria in section 00.

Dashboard (Looker Studio, always on)

Sources: GA4 (events and conversions), Search Console (queries, pages), a Google Sheet or BigQuery table fed nightly from /data (rankings by cluster, CWV, uptime, job status), and GoHighLevel pipeline exports (leads by stage). Pages: Overview, Leads and Attribution, Organic by Cluster, Local, Technical Health, Content Pipeline. The dashboard link is in every digest.

28 / TECHNICAL STACK AND REPOSITORY LAYOUT

Static first, typed content, no page builder

Volume 1 allowed Claude Code to choose the framework. This section makes the recommendation concrete so the automation layer has a stable target.

Recommended stack

Concern	Choice	Why
Framework	Next.js (App Router) with static generation for all marketing routes; ISR only if a route genuinely needs it	Server-rendered HTML for crawlers, strong image/font tooling, first-class Vercel previews, huge maintainer pool
Alternative	Astro (if the owner prefers near-zero JS by default)	Equivalent SEO outcome; slightly less ecosystem for forms/serverless. Either is acceptable; do not use WordPress/Elementor for this property - the automation layer depends on content in Git.
Language	TypeScript, strict mode	Content schemas are types; broken content fails the build, not production
Styling	Tailwind CSS with design tokens defined once in tailwind.config (section 29)	Small CSS, consistent tokens, no runtime styling library
Content	MDX for Insights/case studies; JSON/TS objects for nav, services, industries, FAQs, proof, testimonials, SEO metadata	Diffable, reviewable in PRs, lintable, no CMS dependency
Forms	Route handler (serverless) that validates server-side, checks anti-spam (Turnstile or hCaptcha invisible mode), forwards to GoHighLevel inbound webhook, and returns the thank-you state	Keeps secrets server-side; deterministic event firing
Hosting	Vercel (Pro) or Cloudflare Pages + Workers	Preview deploys per PR, atomic deploys, rollback, edge caching
Analytics	GTM (web container) + GA4 + Google Ads tags with Consent Mode v2; optional server-side GTM later	Section 32
Booking	GoHighLevel calendar (preferred, keeps data in one CRM) or Calendly; loaded on interaction	Attribution continuity
Images	next/image with AVIF/WebP, explicit dimensions, hero preloaded	CWV
Fonts	Two self-hosted WOFF2 files max (one headline weight, one body weight) via next/font, with a system fallback stack	CWV, no layout shift
Search provider	Adapter: Semrush API / DataForSEO / SerpApi [OWNER picks]	Section 19-20
Uptime	Adapter: Checkly / Better Uptime / UptimeRobot [OWNER picks]	Section 24

Repository layout

```
bold-clicks/
  CLAUDE.md          # operating instructions for Claude Code (section 36)
  README.md          # setup, deploy, rotate secrets, run jobs locally
  app/
    (marketing)/
      page.tsx        # /
      services/page.tsx
      google-ads-management/page.tsx  ... other platform pages
      conversion-tracking/page.tsx
      industries/[slug]/page.tsx
      case-studies/[slug]/page.tsx
      insights/[slug]/page.tsx
      austin-ppc-agency/page.tsx
      about/page.tsx  contact/page.tsx  privacy/page.tsx  terms/page.tsx
    (landing)/
      free-ad-audit/page.tsx
      lp/[slug]/page.tsx      # paid campaign variants (noindex by default)
      thank-you/page.tsx      # noindex
    api/lead/route.ts        # form handler -> GoHighLevel
    sitemap.ts  robots.ts   # generated from the canonical route registry
  components/
    # Hero, ProofStrip, ServiceCard, CaseStudyCard, ProcessSteps,
    # CTASection, FAQAccordion, LeadForm, Breadcrumbs, AuthorCard,
    # FirstHand, Calculator, StickyMobileCTA

  content/
    nav.ts  services.ts  industries.ts  faqs.ts  proof.json  testimonials.json
    citations.json  seo.ts (metadata registry)  routes.ts (canonical route registry)
    insights/*.mdx  case-studies/*.mdx
  lib/
    seo.ts (metadata + JSON-LD helpers)  analytics.ts (dataLayer helpers)
    schema.ts (Zod content schemas)  attribution.ts (UTM/click-id capture)
  automation/
    providers/  collect/  analyze/  generate/  report/  config.ts
    lint/no-fabrication.ts
  data/
    keywords.json  approved-facts.json  serp/  gsc/  ga4/  vitals/  runs/
  docs/
    voice.md  runbooks/  digests/  competitive/  decisions.md
  infra/
    dns.json  redirects.json  headers.json
  tests/
    e2e/ (Playwright)  unit/  lighthouse/  a11y/
  .github/workflows/      # ci.yml + one workflow per job in section 25
```

Content schemas (Zod, enforced at build)

Service	{ slug, title, h1, metaTitle, metaDescription, intro, forWho[], deliverables[], measurement[], process[], faqs[], proofRefs[], relatedIndustries[], cta }
Industry	{ slug, title, metaTitle, metaDescription, audienceBehavior, campaignStructure, wastePatterns[], kpis[], objections[], faqs[], proofRefs[], firstHand, publishable: boolean }
CaseStudy	{ slug, clientProfile {industry, geography, model, nameWithPermission?}, startingPoint, objective, changes[], outcome {metrics[{name, value, definition, dateRange}]}, context, permissionOnFile: boolean, verifiedBy, verifiedAt }
Insight	{ slug, title, metaTitle, metaDescription, author, reviewedBy, reviewedAt, publishedAt, updatedAt?, materialUpdate?: boolean, targetQuery, supportsRoute, factIds[] }
Proof	{ id, type: logo testimonial badge metric, source, permissionOnFile, expiresAt?, value }
Fact	{ id, statement, value, unit, dateRange, source: 'accounts' 'public', approvedBy, approvedAt }
Location	{ slug, metro, marketCommentary, localProofRefs[], localFaqs[], serviceRadius, publishable }

Environment variables and secrets (document each in README)

Name	Used by	Where stored
GHL_WEBHOOK_URL, GHL_API_KEY, GHL_LOCATION_ID	Lead route handler, case-study trigger, digest sender	Hosting env + Actions secrets
GTM_CONTAINER_ID, GA4_MEASUREMENT_ID, GA4_PROPERTY_ID, GA4_API_SECRET	Layout, analytics collector, synthetic form test	Hosting env (public IDs) / Actions secrets
GOOGLE_SERVICE_ACCOUNT_JSON	Search Console, GA4 Data API, Indexing checks	Actions secret (base64)
PSI_API_KEY	Vitals collector	Actions secret
SERP_PROVIDER, SERP_API_KEY	Competitive + rank collectors	Actions secret
UPTIME_PROVIDER, UPTIME_API_KEY	Uptime summary collector	Actions secret
TURNSTILE_SECRET (or HCAPTCHA_SECRET)	Lead route handler	Hosting env
ANTHROPIC_API_KEY	Claude Code GitHub Action	Actions secret
NEXT_PUBLIC_SITE_URL, NOINDEX (true on staging)	Metadata, robots, canonical	Hosting env per environment

Security headers (infra/headers.json)

- Strict-Transport-Security with preload; Content-Security-Policy that allows only self, GTM/GA/Ads, the booking provider, and the anti-spam provider (nonce-based for inline GTM bootstrap); X-Content-Type-Options nosniff; Referrer-Policy strict-origin-when-cross-origin; Permissions-Policy limiting camera/microphone/geolocation; X-Frame-Options SAMEORIGIN (except the landing layout if it must be embedded - it should not).
- Staging: Basic Auth or Vercel password protection plus NOINDEX=true plus X-Robots-Tag: noindex. The daily-integrity job fails if NOINDEX is true on production.

29 / DESIGN SYSTEM TOKENS AND COMPONENTS

Same DNA as Stoneridge, sharper, darker

These tokens implement Volume 1 section 02. Define them once; every component reads from them.

Color tokens and usage rules

Token	Hex	Use	Never
--bc-black	#0B0D0C	Page background for hero and CTA sections, headline text on light sections, primary cards on light backgrounds	As body-text color on forest
--bc-forest	#2F4428	Primary button on light/cream, section labels, links on cream, icon fills	Large text blocks (contrast with black is poor)
--bc-sage	#6C815F	Dividers, chart secondary series, hover states, subtle backgrounds at 8-12% opacity	Body text; small text under 14px
--bc-cream	#F3F2E7	Default light section background, button text on black, card surfaces	Text on white
--bc-gold	#C6A15B	1-2px rules, eyebrow labels, the cursor/click micro-accent, active nav indicator, focus ring accent (paired with a dark outline)	Buttons, body text, backgrounds
--bc-ink	#1C1F1D	Body text on light backgrounds (softer than pure black)	-
--bc-mist	#C9CCC6	Secondary text on black backgrounds	Primary text

Contrast: every text/background pair must pass WCAG AA (4.5:1 body, 3:1 large text). Cream on black, ink on cream, and cream on forest all pass. Gold on black passes only at large sizes; gold on cream does not pass for text - use it for non-text accents only.

Typography scale

Role	Size / line-height (desktop)	Mobile	Weight
Display (H1 hero)	64 / 1.02	40 / 1.05	Headline font, bold
H1 (interior)	44 / 1.08	32 / 1.1	Bold
H2	32 / 1.15	26 / 1.2	Bold
H3	22 / 1.25	20 / 1.3	Semibold
Body	18 / 1.6	17 / 1.6	Regular
Body small / captions	15 / 1.5	14 / 1.5	Regular
Eyebrow	13 / 1.2, letter-spacing 0.12em, uppercase	12	Bold, gold or sage
Button	16 / 1	16	Semibold

Font choice: one geometric/grotesk headline family and one highly legible body family, both self-hosted WOFF2, or a single variable font used at two weights. Fallback stack: system-ui, -apple-system, Segoe UI, Roboto, Helvetica, Arial, sans-serif. Load headline weight with preload; body weight swap.

Spacing, layout, motion

- Spacing scale: 4, 8, 12, 16, 24, 32, 48, 64, 96, 128 px. Section padding 96px desktop / 64px mobile. Container max width 1200px; text measure max 68ch.
- Breakpoints: 640, 768, 1024, 1280. Mobile first. Grid: 12 columns desktop, 4 mobile.
- Corners: 6px on cards and buttons; 0 on full-bleed sections. Shadows: one subtle elevation for cards on cream; none on black.
- Motion: 150-250ms ease-out for hover lift (2px) and reveal (opacity + 8px translate) on first viewport entry only; honor prefers-reduced-motion by disabling transforms. No parallax, no autoplay video, no continuous animation.
- Ridge motif: a single thin angled line or a cropped ridge silhouette used in the hero and the final CTA only. Never repeat the logo mountain as decoration.

Component specifications

Component	Content and behavior	Rules
Header	Logo (approved asset, cream panel on black or on cream), 5 nav items, primary CTA button; collapses to logo + CTA + menu button under 1024px	Sticky; shrinks on scroll; CTA always visible; keyboard operable; aria-expanded on menu
Hero	Eyebrow, H1, subhead, primary + secondary CTA, optional single proof line; black background	LCP element is the H1 text (no hero image required); if an image is used, it is preloaded and sized
ProofStrip	Row of proof items from proof.json with permissionOnFile=true only; renders process proof if empty	Never renders placeholders; hidden entirely if no verified proof
ServiceCard	Icon, title, 2-line description, link	Whole card clickable with a single accessible link; hover lift
ProcessSteps	5 steps (Audit, Build/Rebuild, Launch, Optimize, Scale) each with 'what you get'	Ordered list semantics
CaseStudyCard	Client type, starting problem, action, verified outcome with metric definition and date range, link	Renders only when permissionOnFile and verifiedAt exist
IndustryGrid	Cards for industries with publishable=true	Unpublishable industries do not appear anywhere
MeasurementBlock	Diagram of lead -> qualified -> booked -> closed with tracking layer explanation	Inline SVG, accessible description
FAQAccordion	Visible questions/answers; expanded by default on mobile for the first item	Native details/summary or button + region with aria-controls; content in HTML regardless of state (crawlable)
CTASection	Black section, headline, one-line supporting copy, form or button	Fires cta_click with location attribute
LeadForm	Fields per section 31; inline validation; hidden attribution fields; submit state; error summary	Labels always visible; autocomplete attributes; error messages linked with aria-describedby; no CAPTCHA challenge visible unless risk score demands
StickyMobileCTA	Appears after scrolling past hero on mobile; single button	Does not cover content; respects safe-area insets; dismissible; not shown on the audit page (form is already there)
Breadcrumbs	Home > Section > Page	BreadcrumbList JSON-LD generated from the same data
AuthorCard / FirstHand	Named operator with role and short bio; FirstHand renders a visually distinct block with a 'From our accounts' label	FirstHand is required on Insights (linter)
Calculator	Break-even CPL/CAC: inputs (close rate, average revenue, margin, target ROAS), outputs (max CPL, CAC), methodology link	Works without JS for the explanation; results computed client-side; accessible form controls; fires pricing_or_qualifier_interaction
Footer	Legal name, address (if public), phone, email, nav, legal links, social links, Stoneridge family line	Real details only [OWNER]

30 / PAGE COPY DRAFTS

Working copy, written to the voice rules

Implement as-is, then let the owner edit in the content files. Anything in [brackets] is an owner input, not copy.

Homepage

Section	Draft copy
Eyebrow	PAID MEDIA, BUILT TO PERFORM
H1	We run ads. That's it.
Subhead	Google, Meta, Microsoft and YouTube campaigns built around qualified leads, booked calls and revenue - with tracking that shows what actually worked.
CTAs	Primary: Get My Free Ad Audit. Secondary: See How We Work.
Proof line (only if verified)	[Managing \$X/month across Y accounts] or, pre-proof: 'Specialists in paid media for healthcare, legal and home-services businesses across Texas.'
The problem (H2)	Most ad accounts leak money in the same five places.
The problem (body)	Search terms nobody reviews. Conversions that count every dial as a lead. Landing pages that don't match the ad. Budgets spread across campaigns that never had a chance. Reports that stop at clicks. We fix those first, because that's where the money is.
Section (continued)	Draft copy
What we manage (H2)	Four platforms. One job: leads that close.
Cards	Google Ads - Search and Performance Max built around your highest-value services. / Meta Ads - Creative testing and audiences that produce leads, not likes. / Microsoft Ads - Cheaper clicks from an audience most agencies ignore. / YouTube Ads - Video that fills the top of the funnel and gets measured like everything else. / Conversion Tracking - Calls, forms, CRM outcomes and offline conversions wired back to the platforms.
Why Bold Clicks (H2)	Specialists, not generalists.
Why (bullets)	Ads are the only thing we do, so we get very good at them. / You own your ad accounts and your data - always. / Every call, form and booked meeting is tracked to its source. / We report on qualified leads and revenue, not impressions. / One accountable operator, direct access, no account-manager layer.
How it works (H2)	Audit. Build. Launch. Optimize. Scale.
Steps	Audit - We review your accounts, tracking and landing pages and show you exactly where spend is leaking. / Build or rebuild - Campaign structure, conversion tracking, negatives, creative and landing pages set up right. / Launch - Live within [X] days of kickoff with tracking verified before the first dollar is spent. / Optimize - Weekly search-term, bid, creative and budget work tied to lead quality feedback from your CRM. / Scale - When cost per qualified lead is proven, we add budget and platforms deliberately.
Case studies (H2)	Proof, with the numbers defined.
Industries (H2)	Built for businesses where a lead is worth real money.
Measurement (H2)	A lead is not the finish line.
Measurement (body)	We track the click, the call, the form, the booked appointment and - where your CRM allows - the closed customer. Then we send that back to Google and Meta so the algorithms optimize for people who actually buy.

Section (continued)	Draft copy
FAQ (H2)	Straight answers.
Final CTA (H2)	Want to know where your ad spend is leaking?
Final CTA (body)	Request a free audit. You'll get a written review of your accounts, tracking and landing pages - whether or not you hire us.

Homepage FAQ (visible; answers are owner-editable)

Question	Draft answer
What's the minimum ad spend you work with?	[OWNER: state the real minimum] We're honest when a budget is too small to produce reliable data, and we'll tell you what to do instead.
Who owns the ad accounts?	You do. We work inside accounts you own, and if we ever part ways you keep everything - campaigns, history, data.
How do you charge?	[OWNER: flat monthly management fee / percentage / hybrid]. No long-term contracts; month to month after the initial [X]-day build period.
How fast can campaigns launch?	Typically [X] days from kickoff, and never before conversion tracking is verified.
Which platforms do you manage?	Google Ads (Search, Performance Max, Shopping where relevant), Meta (Facebook and Instagram), Microsoft Ads and YouTube. We don't do SEO, web design or social management - our sister agency Stoneridge Digital does.
What will I see in reports?	Leads, qualified leads, booked calls, cost per qualified lead, and revenue where your CRM allows - by platform and campaign. Plus what we changed and why.
Do you work with businesses outside Austin?	Yes. We're based in Austin and work with businesses across Texas and neighboring states, and remotely nationwide.

Free Ad Audit page (/free-ad-audit/)

Element	Draft
Layout	Landing layout: logo, no nav, form above the fold on desktop, form directly after the headline block on mobile
Title tag	Free PPC / Ad Account Audit Bold Clicks
H1	Find out where your ad spend is leaking.
Subhead	A written review of your Google, Meta, Microsoft or YouTube accounts - campaign structure, tracking, search terms, creative and landing pages - delivered in [X] business days. No pitch deck. No obligation.
What you get (list)	A score for each of the five leak points. / The specific search terms, audiences or placements wasting budget. / Whether your conversion tracking counts real leads or noise. / Landing-page fixes ranked by impact. / A plain-English plan you can act on with or without us.
Who it's for	Businesses spending [OWNER: real threshold] or more per month on ads, or ready to start. Healthcare, legal, home services and other lead-driven businesses in Texas and beyond.
How it works	1. Request the audit. 2. Grant read-only access (we'll show you how - takes two minutes). 3. Get your written audit and a 30-minute walkthrough call.
Trust	You own your accounts. Read-only access only. We never make changes during an audit. [Named operator] reviews every account personally.
Form CTA	Get My Free Ad Audit
Thank-you	Got it. [Name], we'll review [company] and be in touch within one business day. Want to lock in your walkthrough call now? [Calendar]. Fires generate_lead exactly once.

Google Ads Management page (/google-ads-management/) - structure and key copy

- **Title:** Google Ads Management | Bold Clicks. **H1:** Google Ads management that's measured in qualified leads, not clicks.
- **Above the fold:** For lead-driven businesses spending [threshold]+ per month. We manage Search, Performance Max and Shopping campaigns, own the tracking, and report on what booked and what closed. CTA: Get My Free Ad Audit.
- **Sections (H2s):** What's included / How we structure accounts / Search vs Performance Max: when we use each / Negatives and search-term hygiene / Tracking that counts real leads / Landing-page alignment / What the first 30 days look like / Reporting / Case studies / FAQ / CTA.
- **First-hand block:** [OWNER/operator: one real account observation, e.g., a search-term waste pattern found in a Texas home-services account and what changed].
- **Internal links:** /conversion-tracking/, /industries/home-services/, /industries/legal/, /austin-ppc-agency/, the PMax vs Search Insight, the audit checklist Insight.
- **FAQ topics:** account ownership, budget minimums, PMax risk, how long until results, what happens to existing campaigns, Local Services Ads.

Austin page (/austin-ppc-agency/)

Element	Draft
Title tag	Austin PPC Agency - Google & Meta Ads Management Bold Clicks
H1	Austin's paid-media specialists.
Intro	Bold Clicks is an Austin-based agency that does one thing: run ads that produce qualified leads for local businesses. We work from [neighborhood/area] with clients across the metro - Round Rock, Cedar Park, Georgetown, Pflugerville, Leander, Lakeway and San Marcos - and across Texas.
H2: What advertising in Austin actually looks like	Austin is one of the most competitive local ad markets in the country for home services, med spas, dental and legal. [FirstHand: real CPC/competition observations from Bold Clicks or Stoneridge accounts, with date and vertical. Until available: qualitative observations only.]
H2: Industries we serve in Austin	Cards linking to publishable industry pages, each with one Austin-specific sentence.
H2: Austin case studies	CaseStudyCards filtered to Austin metro; if none yet, render 'Austin account lessons' from approved first-hand facts instead - never an empty section, never fake cards.
H2: Local FAQs	Minimum spend for Austin campaigns; whether Local Services Ads make sense for Austin home services; how service-radius targeting works for the suburbs; whether we meet in person.
Local trust	Real address (if public) with embedded map, phone, hours, GBP link, Austin Chamber membership badge (only if real), named operator.
Schema	Organization/ProfessionalService with address only if public; areaServed listing Austin metro cities; sameAs to GBP and profiles.

About page outline

- H1: The paid-media arm of the Stoneridge family.
- Sections: Why we only run ads / Who does the work (named operators with real bios and photos) / Our relationship to Stoneridge Digital and Stoneridge Equity (accurate description [OWNER]) / How we work with clients (ownership, access, reporting) / Platform experience (only claims from proof.json) / Where we are (Austin, Texas) / CTA.
- No stock photos, no invented team members, no 'award-winning' without the award in proof.json.

31 / LEAD CAPTURE, ROUTING AND CRM

Every lead arrives in GoHighLevel with its source attached

The form is short. The data model behind it is complete. Nothing is lost between the click and the pipeline.

Form fields

Field	Type	Required	Validation / notes
Full name	text, autocomplete=name	Yes	2-80 chars
Business email	email, autocomplete=email	Yes	Syntax check; warn (not block) on free-mail domains
Company or website	text, autocomplete=organization	Yes	Normalize URL if it looks like one
What do you need help with?	select	Yes	Google Ads / Meta Ads / Microsoft Ads / YouTube Ads / Tracking & attribution / Not sure - audit everything
Monthly ad spend (approx.)	select	No (recommended)	Not running ads yet / Under \$2k / \$2k-5k / \$5k-15k / \$15k-50k / \$50k+ [OWNER may adjust bands]
Industry	select	No	Dental / Ortho / Med spa / Healthcare (other) / Legal / Home services / Other
Phone	tel, autocomplete=tel	No	E.164 normalization; SMS consent checkbox appears only if phone entered
Consent	checkbox	Yes (where required)	Privacy policy link; SMS consent separate and unchecked by default

Hidden attribution fields (captured client-side, validated server-side)

On first visit, lib/attribution.ts stores first-touch values in a first-party cookie (365 days) and session values in sessionStorage: utm_source, utm_medium, utm_campaign, utm_content, utm_term, gclid, gbraid, wbraid, fbclid, msclkid, landing_page, referrer, first_touch_timestamp. On submit, both first-touch and last-touch sets are attached, plus page_url, device type, GA4 client_id (read from the _ga cookie), and a server-generated lead_id.

Webhook payload to GoHighLevel (inbound webhook, JSON)

```
{
  "lead_id": "bc_01J...",
  "submitted_at": "2026-09-18T15:04:05-05:00",
  "form": "free-ad-audit",
  "page_url": "https://bold-clicks.com/free-ad-audit/",
  "name": "...", "email": "...", "phone": "+1512...", "company": "...",
  "need": "google_ads", "spend_band": "5k-15k", "industry": "home_services",
  "consent_privacy": true, "consent_sms": false,
  "first_touch": { "source": "google", "medium": "cpc", "campaign": "...", "content": "...", "term": "...",
    "gclid": "...", "landing_page": "...", "referrer": "...", "timestamp": "..." },
  "last_touch": { "source": "google", "medium": "organic", "campaign": "gbp", "landing_page": "...",
    "timestamp": "..." },
  "ga_client_id": "GA1.1.123.456", "device": "mobile", "geo_hint": "Austin, TX (from request headers)"
}
```

Map every field to a GoHighLevel custom field. Tag the contact with the form name, need, and spend band. Trigger the 'New website lead' workflow: instant SMS + email to the owner, autoresponder to the lead with the calendar link, task created for follow-up within 1 business hour, and pipeline stage 'New'.

Routing and speed-to-lead

- Spend band \$15k+ or industries Legal/Med spa: flag 'priority' and send an additional SMS to the owner.
- 'Not running ads yet' with spend under \$2k: autoresponder includes the budget-planning guide and calculator; still enters pipeline, lower priority.
- Owner SLA: first human response within 1 business hour, 9am-6pm CT. GHL workflow escalates to Jay [OWNER confirm] if untouched after 2 hours.
- Pipeline stages must include Qualified and Closed so automation can compute qualified-lead and close attribution for the monthly report and feed offline conversions (section 32).

Anti-spam and failure handling

- Honeypot field + Turnstile/hCaptcha invisible mode + rate limiting by IP. No visible challenge unless the risk score demands it.
- If the GHL webhook fails, the route handler queues the payload (hosting KV/queue) and retries with backoff for 24 hours, emails the raw lead to the owner immediately, and the synthetic test alert fires. The visitor still sees the thank-you state.
- All payloads are logged (redacted) for 30 days so a lost lead can always be recovered.

32 / ANALYTICS, TAG MANAGER AND ATTRIBUTION

Track outcomes, feed them back, respect consent

Implements Volume 1 section 11 concretely.

Data layer event schema

```
dataLayer.push({ event: 'cta_click', cta_location: 'hero|header|sticky|section-|footer', cta_text: '...' });
dataLayer.push({ event: 'form_start', form_id: 'free-ad-audit' });
dataLayer.push({ event: 'generate_lead', form_id: 'free-ad-audit', lead_id: 'bc...', need: 'google_ads',
  spend_band: '5k-15k', value: 0, currency: 'USD' });
dataLayer.push({ event: 'phone_click', phone_location: 'header|footer|austin-page' });
dataLayer.push({ event: 'calendar_open', provider: 'ghl|calendly' });
dataLayer.push({ event: 'appointment_booked', lead_id: 'bc...' }); // from booking callback or thank-you param
dataLayer.push({ event: 'case_study_view', case_study: '' });
dataLayer.push({ event: 'pricing_or_qualifier_interaction', tool: 'calculator|spend_band', value_bucket: '...'
});
```

GTM container blueprint

Tag	Trigger	Notes
GA4 Configuration (Google tag)	Consent initialization -> All pages	Consent Mode v2 defaults set before the tag loads; server_container_url if server-side GTM is added later
GA4 Event: each data-layer event above	Custom event trigger matching the event name	Mark generate_lead and appointment_booked as key events in GA4
Google Ads Conversion: Lead	generate_lead	Enhanced conversions enabled with hashed email/phone from the data layer, only when ad_user_data consent granted
Google Ads Conversion: Booked appointment	appointment_booked	Secondary conversion; optimize to Lead initially
Google Ads Remarketing	All pages	Consent-gated
Meta Pixel + CAPI (via GHL or server-side)	PageView; Lead on generate_lead	Deduplicate with event_id = lead_id
Microsoft UET	All pages; Lead goal on generate_lead	Consent-gated
Call tracking (optional: CallRail or GHL numbers)	All pages	Dynamic number insertion only if a call-tracking vendor is chosen; count calls over 60s or tagged qualified as conversions
Consent management platform	First	Cookie banner appropriate to jurisdictions [OWNER: primarily US; add GDPR handling if EU traffic is meaningful]

Offline conversion feedback loop

1. GoHighLevel pipeline stage changes (Qualified, Closed) fire a GHL webhook to an automation endpoint with lead_id, gclid/wbraid/gbraid, email/phone (hashed), stage, timestamp, and deal value if known.

2. Automation writes the row to data/offline-conversions/ (private) and uploads to Google Ads via the Data Manager / enhanced conversions for leads route (API or scheduled sheet upload) and to Meta via CAPI. Follow current Google documentation at build time; the exact upload mechanism changes.
3. The monthly report shows lead -> qualified -> closed rates by campaign so the owner can see which campaigns produce customers, not just forms.
4. If consent for ad_user_data was not granted for that lead, skip the upload and record why.

Verification (built into CI and the daily-integrity job)

- Playwright test with GA4 DebugView enabled asserts each event fires once with the right parameters on the right page.
- GTM container is exported to infra/gtm-container.json after every published change (via GTM API); the daily job diffs live vs repo and flags drift.
- Consent test: with consent denied, assert no _ga, _gcl, _fbp, or _uetsid cookies before interaction.

33 / CI/QA PIPELINE AND PRE-LAUNCH CRAWL

If it matters, it is a failing check

Required checks on every pull request. A PR cannot merge (and therefore cannot publish) until all pass.

Check	Tool	Passing condition
Lint + typecheck	ESLint, tsc --noEmit	Zero errors
Content schema validation	Zod schemas over /content	All content parses; unpublishable items excluded from routes
No-fabrication lint	automation/lint/no-fabrication.ts	Zero blocks (section 26)
Build	next build	Succeeds; route registry matches generated pages
Metadata audit	Script over build output	Unique titles/descriptions/H1; canonical + OG on every indexable page; noindex on landing variants, thank-you, staging
Structured data	Validator over JSON-LD	Zero errors; Organization on home; BreadcrumbList on hierarchical pages; Article on Insights; no FAQPage rich-result assumptions
Lighthouse CI (mobile)	lhci with budgets	Performance >= 90, Accessibility >= 95, Best Practices >= 95, SEO = 100 on /, /free-ad-audit/, one service page, one industry page, one Insight
Performance budgets	lhci budgets.json	JS <= 150 KB gzipped on marketing routes; total page weight <= 600 KB above fold; LCP <= 2.0s lab mobile (headroom for 2.5s field)
Accessibility	axe-core via Playwright	Zero serious/critical violations; keyboard traversal test of header, form, accordion, mobile menu
Link check	Crawler over the built site	Zero internal 4xx/5xx; external links checked weekly, not per PR
Redirects	Test infra/redirects.json	Every rule resolves in one hop with 301
E2E	Playwright	Form submit -> thank-you -> generate_lead; phone link; calendar loads; sticky CTA behavior; reduced-motion respected
Visual regression	Playwright screenshots (Percy/Chromatic optional)	No unapproved diffs on key pages at 375px and 1280px
Preview deploy	Vercel	Preview URL posted to the PR (used in the weekly digest)

Pre-launch crawl (run once before DNS cutover and monthly afterward)

- Every canonical URL in content/routes.ts returns 200 on production and appears in sitemap.xml exactly once.
- No URL outside the registry is reachable through internal links (no orphan or accidental routes).
- robots.txt allows all public assets; disallows /api/, /thank-you/, /lp/ (or lp variants carry noindex meta).
- X-Robots-Tag and meta robots agree with the intended index state on every URL.
- HTTP -> HTTPS, www <-> non-www (one chosen host), trailing-slash policy: each resolves in a single 301.
- 404 page returns a real 404 status, not 200.
- Hero H1 and primary copy present in initial HTML (view-source test), not injected by JS.
- Images: explicit width/height, modern formats, below-fold lazy, hero not lazy.

- Fonts: two files max, preload on headline weight, font-display swap on body.
- Third-party scripts: inventory documented in docs/decisions.md with load strategy for each.
- Search Console property verified (DNS TXT), sitemap submitted, and URL Inspection shows 'URL is on Google' for the homepage within the first week.

34 / LAUNCH PLAN AND POST-LAUNCH 30/60/90

Eight weeks to launch. Then the system takes over.

Assumes one Claude Code-driven build with the owner supplying inputs from section 35 in week 1.

Week	Build	Owner	Exit criteria
1	Repo, CLAUDE.md, design tokens, component library, route registry, content schemas, CI skeleton	Deliver section 35 inputs; approve font choice and homepage copy	Storybook-style component preview approved on mobile
2	Homepage + free-audit page + lead route handler + GHL webhook + thank-you; attribution capture	Confirm GHL fields/workflow; test a lead end to end	Synthetic form test green; Lighthouse budgets met
3	Services hub + 5 platform pages + conversion-tracking page; FAQ system; breadcrumbs; metadata + JSON-LD helpers	Provide first-hand blocks for each service page	No-fabrication lint green with real first-hand content
4	Austin page; About; Contact; legal pages; industry template + publishable industry pages (home services, legal, med spa, dental at minimum)	Confirm which industries are publishable; supply proof.json entries	Doorway test passes on Austin page
5	GTM container, GA4, Ads conversions, Consent Mode, call/booking tracking; offline conversion pipeline; dashboard	Provide account IDs; grant GA4/Ads/GSC access to service account	Events verified in DebugView; consent test green
6	Automation: collectors, analyzers, no-fabrication linter, weekly digest, synthetic tests, integrity job, alert routing via GHL	Receive first test digest; confirm SMS routing	All section 25 jobs run green on schedule for 3 days on staging
7	Case-study and Insight templates; first two Insights (calendar month 1); calculator; GBP setup	Review and approve first two Insights; verify GBP	Drafts reviewed by a named human; GBP verified
8	Pre-launch crawl; accessibility pass; performance pass; runbooks; README; DNS cutover plan; backup	Approve launch; supply DNS access or apply records	Section 33 crawl checklist complete; rollback plan documented

DNS cutover procedure

1. Lower TTL on existing records to 300s at least 24 hours before cutover (if any records exist).
2. Verify staging is password-protected and NOINDEX=true; verify production env has NOINDEX=false.
3. Add the hosting provider's A/CNAME records; confirm certificate issuance; test HTTPS on both host variants.
4. Flip production; run daily-integrity job manually; run the pre-launch crawl against the live domain.
5. Submit sitemap in Search Console; request indexing for the homepage, audit page, and Austin page.
6. Publish GBP website link with UTM; publish social profile links.
7. Watch uptime and tracking for 24 hours; send the first real weekly digest the following Monday.

Post-launch 30 / 60 / 90

Window	Focus	Success signal
Days 1-30	Indexing coverage, tracking accuracy, first leads, GBP reviews from existing Stoneridge-family clients (only where appropriate and real), citations built, first 4 Insights live, first case study if data allows	All canonical pages indexed; at least one organic lead; digest running clean 4 weeks straight
Days 31-60	Decay/opportunity loop producing its first refresh PRs; Austin map-pack tracking; competitive scan #2; paid campaigns for Bold Clicks itself using the /lp/ template (the agency should run its own ads well)	Austin cluster terms entering top 20; CWV field data appearing in CrUX; cost per qualified lead known for Bold Clicks' own campaigns
Days 61-90	Second and third case studies; first benchmark asset if sample allows; link outreach on the calculator and audit checklist; Texas hub page assessment against unlock criteria	Tier 1 unlock criteria in section 00 met or a written plan for the gap

35 / OWNER-SUPPLIED INPUTS CHECKLIST

What Claude Code cannot invent

Deliver these in week 1. Anything missing becomes a labeled placeholder and is listed in the hand-off report; the site launches with those sections hidden, never faked.

Brand and legal

- Approved logo files: SVG master, PNG on cream, PNG on black, square mark for favicon/social - exactly as supplied, no redraws.
- Legal entity name for footer/terms/privacy; registered address; whether a public office address exists for Austin (and if so, the exact address); public phone; public email.
- Accurate one-paragraph description of the relationship to Stoneridge Digital and Stoneridge Equity for the About page.
- Named operators: names, roles, headshots, 2-3 sentence bios, LinkedIn URLs, platform experience that can be stated truthfully.

Offer and pricing

- Final primary CTA phrase (default: Get My Free Ad Audit).
- Real minimum monthly ad spend; pricing model; contract terms; typical launch timeline; audit turnaround time.
- Spend-band options for the form; which industries are publishable at launch.

Proof (all with permission on file)

- Client logos cleared for use; testimonials with attribution; any partner/certification badges with expiry dates; any metrics for the proof strip with definitions and date ranges.
- First case study data (or confirmation that none is ready - the site ships with process proof instead).
- First-hand blocks: at least one real account observation per service page and per publishable industry page (bullet points are enough; automation drafts around them).

Accounts and access

- Domain registrar and DNS access (or willingness to apply records); hosting account (Vercel/Cloudflare) [OWNER creates; Claude Code never creates accounts].
- GoHighLevel: location ID, inbound webhook URL, API key, calendar link, custom fields created per section 31, pipeline stages including Qualified and Closed.
- Google: GA4 property + stream, Google Ads account ID + conversion actions, GTM container, Search Console (service-account access), Business Profile ownership.
- Meta Business Manager pixel/dataset ID; Microsoft Ads UET tag ID (if used).
- Choice of SERP/keyword provider and uptime provider, with API keys; anti-spam provider keys; Anthropic API key for the Claude Code Action.
- Alert contacts: owner mobile and email; backup contact (Jay) for escalation.

Content decisions

- Approve docs/voice.md banned-phrase list additions.
- Approve the 12-month calendar or reorder it.
- Set the publishing cap and the benchmark minimum-sample rule.
- Confirm consent/privacy jurisdictions and whether EU traffic matters.

36 / CLAUDE CODE OPERATING INSTRUCTIONS

How the build agent should work in this repository

Put the CLAUDE.md below at the repo root verbatim (adjusting paths if the structure changes). It governs the initial build and every automated run afterward.

CLAUDE.md (repo root)

```
# Bold Clicks - CLAUDE.md

## What this is
Marketing site for Bold Clicks, an ads-only paid-media agency in Austin, TX (Stoneridge family).
Primary conversion: free ad audit. Ranking strategy: Austin -> Texas -> surrounding states -> national.
Two governing documents live in /docs/brief/: Volume 1 (brand, sitemap, templates, launch checklist)
and Volume 2 (ranking, content, automation). Volume 1 wins on brand/no-fabrication; Volume 2 wins on
SEO/automation/implementation.

## Non-negotiables
- Use the approved logo assets in /public/brand exactly. Never redraw, recolor, or generate a logo.
- Never invent numbers, clients, testimonials, badges, awards, addresses, or prices. Use [OWNER]
  placeholders in content files; the linter blocks placeholders from publishing.
- Ads only. Do not describe Bold Clicks as SEO, web design, social, email, or automation agency.
- Static/server-rendered HTML for all primary copy. No copy behind clicks or client-side fetches.
- One location page per real market; never templated city pages. Run the doorway test (Volume 2 s.18).
- Automation drafts and proposes via PRs. It never publishes or changes tracking without a merge.

## Commands
pnpm dev | pnpm build | pnpm lint | pnpm typecheck | pnpm test (unit) | pnpm e2e (Playwright)
pnpm lint:content (schemas + no-fabrication) | pnpm audit:meta | pnpm audit:schema | pnpm audit:links
pnpm lhci | pnpm crawl:prelaunch | pnpm job (runs an automation job locally with .env.local)

## Working rules
- Read /docs/voice.md before writing any copy. Every Insight needs a block.
- Content lives in /content. Do not hard-code copy in components.
- Add every new route to /content/routes.ts (sitemap, robots, canonical, breadcrumbs derive from it).
- Add every new event to /lib/analytics.ts and to docs/tracking-plan.md in the same PR.
- Any change to tracking, consent, redirects, or robots requires a PR labeled 'tracking' or 'seo-critical'
  and a note in docs/decisions.md.
- Prefer small PRs with a preview link and a one-paragraph 'why' for the owner's digest.
- When a required input is missing, do not guess: write the placeholder, add it to
  docs/owner-inputs.md, and continue.

## Definition of done (any task)
Build passes; all CI checks green; content schemas valid; no placeholders in publishable files;
metadata unique; all y and Lighthouse budgets met; e2e green; docs updated; PR description written
for a non-technical reader.
```

Phase prompts (use these to start each build phase)

Phase	Prompt for Claude Code
1 Foundation	Read docs/brief/volume-1.pdf and volume-2.pdf. Create the repo layout in Volume 2 s.28, design tokens in s.29, the component library, content schemas, route registry, and CI skeleton in s.33. Do not write page copy yet. Open a PR with a component preview.
2 Conversion core	Build the homepage and /free-ad-audit/ from Volume 2 s.30 copy and s.31 form spec, including the lead route handler, GoHighLevel webhook, attribution capture (s.31), thank-you state, and data-layer events (s.32). Add the Playwright synthetic form test. Meet Lighthouse budgets on mobile.
3 Service pages	Build /services/ and the five platform pages plus /conversion-tracking/ using the Volume 1 s.05 template and Volume 2 s.20 keyword map. Insert placeholders where owner input is pending and list them in docs/owner-inputs.md.
4 Local + industries + about	Build /austin-ppc-agency/ per s.18 and s.30, the industry template, publishable industry pages, About, Contact, legal pages. Run the doorway test and document results.
5 Measurement	Implement the GTM container blueprint (s.32) as an exportable container JSON, Consent Mode v2, Ads conversions with enhanced conversions, the offline-conversion webhook endpoint, and the verification tests. Document the tracking plan.
6 Automation	Implement every job in s.25 with the provider adapter pattern, thresholds from s.24 in automation/config.ts, run logs, the no-fabrication linter (s.26), the weekly digest (s.27), and alert routing through GoHighLevel. Provide runbooks (s.24).
7 Content	Implement Insight and case-study MDX templates, the calculator, and draft calendar month-1 items with FirstHand blocks left for the owner. Open one PR per piece.
8 Launch	Run the pre-launch crawl (s.33), fix everything it finds, write README and DNS cutover steps, generate the hand-off report listing all outstanding owner inputs and the launch acceptance checklist from Volume 1 s.15 with pass/fail status.

Hand-off report (required output of phase 8)

- Launch acceptance checklist (Volume 1 s.15) with pass/fail and evidence links.
- Outstanding [OWNER] inputs and which sections are hidden until they arrive.
- Job schedule table with the last run status of each job.
- Access inventory: every secret, where it lives, how to rotate it.
- Tier-1 unlock criteria (s.00) and the current status of each.
- First-month digest preview.

37 / RISK REGISTER AND ANTI-PATTERNS

What would quietly ruin this

Each risk lists the control that prevents it. If a control is missing from the build, the build is not done.

Risk	Impact	Control
Fabricated proof slips into copy	Credibility and potential legal exposure; permanent trust loss	No-fabrication linter (s.26); proof.json permission flags; ProofStrip renders nothing rather than placeholders
City-page sprawl to 'rank everywhere'	Doorway-page manual action; sitewide ranking loss	Location page rules and doorway test (s.18); route registry review in PRs
AI content farm drift	Scaled-content policy action; brand becomes generic	First-hand block requirement; publishing cap; named reviewer; banned-phrase lint
Tracking silently breaks	Ads optimize to nothing; leads unattributed; wrong decisions	Synthetic form test every 6h; zero-events check; GTM drift diff; S1 alerts
Accidental noindex on production	Total organic loss	Daily integrity check; CI env assertion; SMS alert
Performance regression from third-party scripts	CWV leaves 'good'; Ads landing-page quality drops	Budgets in CI; script inventory; load-on-interaction rule; weekly PSI
Fake freshness	Trust loss; wasted crawl budget	Lastmod lint; material-update tagging
Owner overwhelmed by alerts	Alerts ignored; real incidents missed	One weekly digest; S1 only for outages/forms/tracking; 24h suppression
Vendor lock-in or API changes	Automation stops working	Provider adapters; run logs flag provider failures; README rotation docs
Single point of failure (one person)	Site unmaintainable	README, runbooks, decisions log, escalation contact
GBP suspension	Map-pack loss	No keyword stuffing in name; real address rules; runbook
Consent violations	Legal and platform policy exposure	Consent Mode v2; weekly consent test

Anti-patterns (reject in code review)

- Any page whose body could be pasted onto another agency's site unchanged.
- Hero sliders, autoplay video, popups before content, chat widgets on first paint.
- A 'Locations' dropdown listing cities without pages that meet the s.18 standard.
- Meta descriptions or titles generated in bulk without the uniqueness audit.
- Numbers in copy without a fact id.
- Publishing-date changes without content changes.
- Tracking changes merged without the 'tracking' label and decisions-log entry.
- Automation that can only be run by hand, or that fails silently.

APPENDIX A

Geography reference

Used by the collectors to classify queries and by content authors for natural local references. Not a list of pages to build.

Area	Places to recognize in queries and content
Austin metro (Tier 1)	Austin, Round Rock, Cedar Park, Georgetown, Pflugerville, Leander, Lakeway, Bee Cave, Dripping Springs, Kyle, Buda, San Marcos, Hutto, Manor, Elgin, Bastrop; Travis, Williamson, Hays, Bastrop counties
Dallas-Fort Worth (Tier 2)	Dallas, Fort Worth, Plano, Frisco, McKinney, Arlington, Irving, Denton, Southlake, Grapevine, Richardson, Garland
Houston (Tier 2)	Houston, The Woodlands, Katy, Sugar Land, Pearland, Cypress, Spring, League City, Conroe, Baytown
San Antonio (Tier 2)	San Antonio, New Braunfels, Boerne, Schertz, Seguin
Other Texas	El Paso, Corpus Christi, Lubbock, Waco, Killeen, Temple, College Station, Tyler, Midland, Odessa, Amarillo, McAllen, Brownsville, Laredo
Oklahoma (Tier 3)	Oklahoma City, Tulsa, Norman, Edmond, Broken Arrow
Louisiana (Tier 3)	New Orleans, Baton Rouge, Lafayette, Shreveport, Lake Charles
Arkansas (Tier 3)	Little Rock, Fayetteville, Bentonville, Rogers, Springdale, Fort Smith
New Mexico (Tier 3)	Albuquerque, Santa Fe, Las Cruces, Rio Rancho

Search Console geo-classification rule

A query is assigned to a tier if it contains any place name above (case-insensitive, with common abbreviations: dfw, okc, nola, atx, htx, satx). Queries with no place name are 'unlocalized' and tracked separately; their localized SERPs are sampled through the SERP provider with an Austin location setting.

APPENDIX B

Event and schema reference

Structured data by route

Route	JSON-LD types	Notes
/	Organization (or ProfessionalService) + WebSite	name, url, logo (approved asset URL), sameAs, contactPoint, address only if public, areaServed
/services/ and platform pages	Service + BreadcrumbList	provider -> Organization; serviceType; areaServed
/industries/*	Service + BreadcrumbList	audience.audienceType
/case-studies/*	Article + BreadcrumbList	author = named operator; datePublished; about
/insights/*	Article or BlogPosting + BreadcrumbList	author, reviewedBy (as contributor), datePublished, dateModified only on material update
/austin-ppc-agency/	LocalBusiness/ProfessionalService (only with a real public address) or Organization with areaServed	geo only if address is public
/free-ad-audit/	WebPage	No Offer markup unless the audit is a genuine free offer with terms stated
FAQs everywhere	None required	Visible FAQs for users; FAQPage schema optional and not expected to produce rich results

GA4 key events and Ads conversion mapping

GA4 event	Key event	Google Ads conversion	Meta / Microsoft
generate_lead	Yes	Lead (primary)	Lead / Lead goal
appointment_booked	Yes	Booked appointment (secondary)	Schedule
phone_click	Yes	Phone click (secondary, low value)	Contact
qualify_lead (from CRM)	Yes	Qualified lead (offline upload)	CAPI custom
close_convert_lead (from CRM)	Yes	Closed customer (offline upload, value)	CAPI Purchase/custom
cta_click, form_start, calendar_open, case_study_view, pricing_or_qualifier_interaction	No	-	-

APPENDIX C

KPI glossary and definitions

Definitions are fixed here so reports, case studies, and benchmarks all mean the same thing.

Term	Definition used on bold-clicks.com
Lead	A form submission, tracked call over the qualifying duration, or booked appointment from a unique contact, deduplicated within 30 days
Qualified lead	A lead the client (or Bold Clicks on the client's behalf) has marked as a fit in the CRM: right service, right area, real intent
Booked call / appointment	A calendar event created by the lead or the intake team, tied to a lead_id
Closed customer	A CRM opportunity marked won, with value when known
Cost per lead (CPL)	Ad spend / leads, for a stated platform, campaign set, and date range
Cost per qualified lead (CPQL)	Ad spend / qualified leads, same scope
CAC	Ad spend (plus management fee if stated) / closed customers
ROAS	Attributed revenue / ad spend; always stated with the attribution model and date range
Wasted spend	Spend on search terms, placements, or audiences that produced zero leads over a stated window and were subsequently excluded
Message match	The ad's promise, keyword/audience, headline, proof, and CTA are consistent on the landing page
Core Web Vitals 'good'	At the 75th percentile of field data: LCP <= 2.5s, INP <= 200ms, CLS <= 0.1
Tier unlock	The measurable conditions in section 00 that permit expanding pages and targeting to the next geography

CLOSING DIRECTIVE FOR CLAUDE CODE

Build bold-clicks.com as specified in Volume 1 and this Volume 2. Sequence the ranking effort Austin first, then Texas, then the surrounding states. Produce fewer, stronger pages with first-hand evidence on every one. Wire every form, call, and booking to GoHighLevel with source data intact and feed qualified and closed outcomes back to the ad platforms. Implement every monitor and scheduled job so the owner receives one weekly digest and is alerted immediately only when revenue-critical systems fail. Draft content and fixes as pull requests; never publish or alter tracking without a merge. Fabricate nothing. When an input is missing, leave a labeled placeholder and keep building.